



MezTal Execution Answers for Gemini (2025-11-23)

1. Content Strategy for Niche Roles (e.g., *Senior Living Billing Analyst, Clinical Recruiter*)

For highly specialized roles with low search volume, the strategy is to **cover them within broader topical hubs** rather than standalone thin pages. We will integrate niche roles into relevant industry or solution content clusters. For example, a role like **Senior Living Billing Analyst** would be featured under a larger *Senior Living/Healthcare Operations* topic, and **Clinical Recruiter** under a *Healthcare HR/Recruitment* hub. This ensures they benefit from the context of a higher-level page and internal linking, avoiding orphan pages with no traffic. Competitor analysis shows Intugo succeeding with *vertical-specific content*, tailoring blog posts to very specific roles in context (e.g. *Nearshore Marketing Manager, Community Management Support*). We'll emulate that approach by **creating content sections or articles that group multiple niche roles together** around a common theme (e.g. a guide to "Key Finance Roles in Senior Living" mentioning Billing Analysts, etc.).

Each niche role will still have a **dedicated CMS entry** (to allow dynamic content insertion and sales reference), but the on-page copy will emphasize how that role fits into the bigger picture. For instance, the *Senior Living Billing Analyst* page can discuss billing challenges in senior living and link to a broader *Healthcare Finance Solutions* page. Likewise, the *Clinical Recruiter* page can be part of a *Clinical Talent Acquisition* series. Internal linking will connect these role pages to their parent hub (pillar page) and to relevant blog articles or resources. This hub-and-spoke model ensures even low-volume roles have SEO value by inheriting authority from the hub and providing highly targeted content for niche queries. It also serves sales enablement: if a prospective client in a niche industry visits the site (or is sent a link by our team), they see content addressing that specific role, lending credibility. In summary, **no niche role is left behind** – each is woven into a content cluster with contextual support, rather than floating alone.

To implement this, we will use the CMS to tag each role with an industry or solution category (e.g., "Senior Living" or "Healthcare Recruiting"). The dynamic role page template can pull in that category's overview content or sidebar links. If certain roles have extremely little to say, we may combine them on one page or mark them as *noindex* (to avoid thin-content penalties), while still keeping the content available for users. However, our preference is to at least publish a short, well-structured page per role, enriched with cross-links, rather than skip them – this caters to the "long tail" of search and signals completeness of our offerings. **In short, niche role pages will be short but authoritative**, bolstered by linking to hub pages and related resources, ensuring Google understands their relevance despite low volume. (*Competitor Intugo's success targeting specific roles reinforces that even niche content can build authority when framed properly.*)

2. Source of Truth & Methodology for the Salary_Range Field in Roles CMS

We plan to make the **Salary_Range** in the Roles CMS a credible, data-driven field that reflects current market rates for each role. The *source of truth* will be a combination of **market salary data** and our **internal placement data**:

- **External Market Data:** We will pull ranges from reputable salary databases (e.g. Glassdoor, Payscale, Robert Half guides) filtered for the role and region (primarily Latin America for nearshore roles, or U.S. if applicable). This provides an up-to-date baseline. For instance, for a *Senior Living Billing Analyst*, we'd find the typical salary range for a billing analyst in healthcare in LatAm (if hiring nearshore) or the U.S. equivalent and then adjust.
- **Internal HubSpot Closed-Won Data:** We will cross-reference our own placement records to fine-tune these ranges. Our HubSpot deal logs show the salary levels clients have agreed to for filled roles (often implied by "Total Estimated Salary" in deals). For example, our data shows we placed a **Billing Analyst** for a senior living client in 2024 at a certain salary (Trustwell Living's deal) – we can use that as a validation point. If our closed-won data consistently shows, say, mid-level accountants in Mexico at \\$X-\\$Y, our CMS range will reflect that. (*Using internal data grounds our content in reality, increasing trust with clients.*)

Methodology: For each role entry, we'll define the salary as a broad range (e.g., "\\$40,000-\\$50,000 USD per year") appropriate to that position's typical compensation. This range will be triangulated by:

1. **Job Level & Experience:** We'll consider if it's entry, mid, or senior level – each has different pay scales.
2. **Location Factor:** We will assume nearshore roles (Mexico/LatAm) – if a role is commonly filled in certain cities (e.g. Guadalajara vs. Mexico City), we'll ensure the range accounts for any regional cost differences, though in many cases differences are minor. (If needed, we might eventually localize salary by location using the Locations CMS, but initially a LatAm-average range is fine.)
3. **Update Frequency:** We will establish a process to update these ranges **quarterly or bi-annually**. The tech job market can shift, so our *source of truth* might be an integrated Google Sheet or JSON that the CMS pulls from, maintained by our team. We could also integrate an API from a salary data provider if available for automation.

Including salary ranges on role pages adds SEO value and answers common user questions ¹. We'll clearly label it (e.g., "Typical Salary Range for a [Role] in LatAm") so Google doesn't mistake it for a job posting salary (more on schema in question 6). This field will be referenced in content (perhaps as part of an FAQ section on the role page: "**What is the salary of a [Role]?**") to capture search queries around compensation ¹.

DATA GAP: We need to finalize the exact data source for these ranges. *Proposed solution:* conduct a one-time analysis using our 2023-2024 placement data exported from HubSpot (to get real salaries placed) and supplement with a salary survey for roles we haven't placed yet. This combined dataset will become our master reference for populating the **Salary_Range** field initially. Then assign an owner (e.g., HR or Content team) to monitor salary trends and adjust the CMS values periodically. By establishing this process, the **Salary_Range field stays accurate and trustworthy**, enhancing our content marketing (it's genuinely useful info) and helping set client expectations.

3. Operational Differentiators – *MezTal vs Competitors (e.g., Toptal)*

When discussing how MezTal stands out in comparison pages (like “MezTal vs Toptal”), we will highlight **operational and service model differentiators** that set us apart from freelance marketplaces and other staffing firms:

- **Full-Service, EOR-Based Staffing vs. Marketplace Model:** Unlike Toptal which is essentially a global freelancer marketplace, MezTal operates as a **hands-on recruiting partner and Employer of Record (EOR)** for nearshore talent ². This means our clients get a *dedicated nearshore team member* with HR onboarding, payroll, and local compliance handled by MezTal (via our FlexTal entity), rather than just a contractor from a platform. The client retains day-to-day control of the worker, but **we handle all back-office employment logistics** – an approach focused on building long-term team capacity, not one-off gigs. Toptal, by contrast, focuses on short-term placements of freelancers worldwide and even offers managed project delivery; MezTal’s differentiator is **permanent team integration**. We emphasize that we **build teams that stick around**, aligning with client culture, whereas a Toptal freelancer might be project-based and more transient ².
- **Latin America & Time-Zone Alignment:** MezTal specializes in **Latin American (nearshore) talent**, meaning our clients work with professionals in (or near) U.S. time zones and with strong bilingual communication. Toptal’s pool is global (150+ countries) ³, which can be a pro for breadth, but a con for time overlap and cultural alignment. We’ll stress that MezTal’s operational model leverages **geographic proximity** – easier collaboration, occasional in-person meetups if needed (e.g., our Guadalajara hub), and generally fewer time-zone headaches. This *focus on LatAm* also means we have deep local networks and know the local market nuances (salaries, work culture) ⁴ ⁵, which a broad platform cannot match.
- **Quality & Retention vs. Speed:** Toptal markets itself on rapid matching (often presenting a candidate in 24-48 hours) and the “top 3%” vetting claim ³. MezTal’s differentiator is **meticulous vetting for long-term fit** rather than instant placement. We highlight our **97% placement success rate and low attrition** (as noted by our internal metrics), showing that we invest more time upfront to find candidates who will stay 2+ years. This contrasts with Toptal’s rapid, project-centric approach – our approach reduces turnover and re-hiring costs for clients. In comparisons, we’ll cite our **180-day guarantee** or similar policies as evidence of confidence in our placements (for instance, if a hire doesn’t work out, we replace them at no cost, demonstrating we prioritize right-fit over quick fill).
- **Personalized Account Management vs. Self-Service:** Clients working with MezTal get a **high-touch experience** – a dedicated account manager/recruiter who learns their business and handles the entire recruitment life-cycle, plus ongoing support (performance check-ins, etc.). With Toptal, while they do have some matching services, the model is more *self-service* once the freelancer is onboarded – the freelancer is independent. MezTal can be positioned as effectively an **extension of the client’s HR team**, providing guidance on role definitions, market salary consulting, and even equipment or workspace for the remote hire ⁴ ⁶. This operational difference means **less burden on the client**. We will include this in comparisons as “white-glove service” vs. “marketplace platform.”
- **Team Building vs. One-off Talent:** MezTal’s model is ideal for companies looking to **build an extended team** (e.g., multiple developers over time, an entire nearshore department). We handle

recruiting at scale and can fill multiple roles in parallel, with consistency in quality. Toptal is typically used to find an individual specialist for a defined task. We'll differentiate by citing **use cases**: e.g., "If you want to *augment your team with several full-time engineers who grow with your company*, MezTal's nearshore staffing is designed for that. If you just need one freelancer for a 3-month project, a marketplace like Toptal might suffice." By clarifying these scenarios, we help prospects choose (and subtly show that for strategic hiring, MezTal adds more value).

- **Cost-Effectiveness & Predictable Pricing:** Operationally, MezTal offers **transparent, predictable pricing** (typically a monthly rate or markup on salary) for full-time hires, often at lower cost than U.S. hires due to the labor arbitrage of nearshore. Toptal talent, being top-tier global, can actually be expensive – plus Toptal adds a premium margin. We will highlight that **nearshore full-time hires through MezTal can save 30-50% vs U.S. salaries**. Also, because they're full-time dedicated staff, their output vs. cost is often better than a part-time expensive freelancer from Toptal. Any **case study or data** (e.g., "Client X saved \\$Y annually by hiring a nearshore team with MezTal instead of via Toptal or local hiring") will reinforce this.

In summary, **MezTal's differentiators** come down to *engagement model* (full-service staffing vs. freelance marketplace), *regional focus* (LatAm expertise), *long-term team emphasis* (retention and cultural fit), and *high-touch service*. Our comparison pages will use these points to position MezTal as the better choice for companies seeking stable, long-term remote team members with minimal administrative hassle. Conversely, we implicitly define when a competitor's model might be for a different need. By doing so, the pages remain fair and useful, but clearly communicate why MezTal provides superior **operational value** in our domain [2](#) [3](#). We will support claims with evidence where possible (testimonials, guarantees, stats) to avoid just platitudes.

4. Programmatic Logic for Linking Roles and Locations Collections (Wix Field Mapping)

We will leverage Wix CMS's reference fields to create a dynamic link between the **Roles** collection and the **Locations** collection, enabling content to be inter-related. The goal is that on a given role's page, we can show relevant locations (e.g., key hubs for that talent), and on a location's page, we can list roles commonly filled there.

Field Mapping Approach: In the **Roles** collection, we'll add a field (likely a multi-reference field) called "**Key Locations**" or similar. This allows us to select one or multiple Location entries that correspond to where that role is significant. For example, for *React Developer* role, we might tag "Guadalajara" and "Mexico City" as key locations (since those cities have strong pools of IT talent). Conversely, in the **Locations** collection, we can have a multi-reference to Roles (e.g., a field "In-Demand Roles in this Location"). However, we might not even need a field on the Locations side if we prefer to manage it one-way – Wix allows us to query the inverse relation. For simplicity, we'll likely manage it from the Roles side only: each Role will have 0, 1, or many associated Locations.

On the site front-end: - Role Pages (/hire/[slug]): Using the reference, we can programmatically pull in a list of linked locations. We might present this as a callout like "Popular in: Guadalajara, Mexico City" with links to those location pages. For instance, the *Python Developer* role page could say "Our strongest Python talent is in Guadalajara" and link to the Guadalajara hub page. This is achieved by adding a repeater or text

element in the role page template bound to the “Key Locations” references. If the field has multiple locations, Wix will automatically provide an array we can loop through. This cross-link not only helps users navigate (maybe a client interested in Python devs will click to see why Guadalajara is highlighted) but also aids SEO by creating a mesh between our role content and location content. - **Location Pages (/locations/[slug])**: We can either manually list top roles (using the multi-ref on Location if we added it), or simply query all Roles that reference that Location. Wix Code (Velo) allows filtering a dataset: e.g., on *Guadalajara* page, we can set up a dataset of Roles filtered where “Key Locations contains Guadalajara”. That would dynamically fetch, say, the top 5-10 roles we have tagged for Guadalajara. We might display that as “Top Talent in Guadalajara: Software Engineers, QA Analysts, Project Managers...” each linking to the respective role page. This gives the location page additional depth, showing what kind of hires one can get there, which reinforces Guadalajara’s position as a talent hub.

From a content management perspective, we’ll maintain these mappings carefully: - Initially, we’ll **bulk assign locations to roles** based on known data (perhaps via the CSV import or in the spreadsheet where we have a column for “Suggested Hub”). For example, our SEO spreadsheet’s *Geographic Hubs* cluster indicates certain cities for certain content (e.g., Monterrey for software dev, etc.), which we can use as a guide. - If a role is broadly available (e.g., *Virtual Assistant* might be common across all LatAm), we might either tag multiple major cities or leave it un-tagged if it’s not location-specific. Un-tagged roles simply won’t show a location emphasis on their page (and that’s fine). - We will also consider using **dynamic tags** on these fields – for instance, a location like “Latin America (General)” for roles that aren’t tied to one city. We could create a Location entry for “Latin America” or “Remote LATAM” to tag roles that are ubiquitous, ensuring the role page still shows something like “Available across Latin America” rather than listing every city.

Wix Implementation Details: Using reference fields will automatically create relationships we can exploit in the Wix Editor design. We’ll design the role page template to include a repeater linked to “Key Locations”; each item in the repeater will display the location’s Name and maybe an image (like a flag or city skyline) and link to the location page (since the reference field gives us the actual Location item, including its URL). No custom code is required for basic linking – Wix’s dataset binding can do it. If more control is needed (say we want to sort locations by a specific order or limit to X items), we can use a bit of Velo code to manipulate the data.

Example: On the “*Data Analyst*” role page, suppose we have tagged “Bogotá” and “Mexico City” as key locations. The page might have a section: “*Where are the best Data Analysts located?*” and display those two cities with a brief note – “*Mexico City – Large fintech sector driving demand*”, “*Bogotá – Strong analytics talent pipeline*”. Each links to the city page. Conversely, on the *Mexico City* page, we’d list Data Analyst among top roles (because it was tagged). This **bi-directional linking** enriches UX and SEO: Google sees that our content is interlinked contextually (roles <> locations), reinforcing relevancy for searches like “Mexico City software developers” or “hire analysts in Latin America”.

Overall, the programmatic logic is to **maintain a many-to-many relationship** between roles and locations in the CMS and use Wix’s dynamic page capabilities to surface that relationship. As we add new roles or new locations, we’ll update the references accordingly. This is scalable – if we launch a new city page, we can quickly assign relevant roles to it, and they’ll start appearing on that city’s page without manual editing (and vice versa). It’s a robust way to implement the *Hub & Spoke* model in our CMS: where “Geographic Hubs” and “Role pages” are interconnected systematically, not just via hard-coded links.

5. Visual Styling Plan for Embedding Workable and LinkedIn Content

When embedding external content from Workable (our jobs board) and LinkedIn, we will ensure the visuals align with the MezTal site's design for a seamless look and feel:

- **Workable Jobs Widget:** We plan to embed Workable's job listing widget on our Careers or Jobs page. The good news is that Workable's embed automatically **inherits our site's CSS** for basic elements like fonts and colors ⁷. This means as soon as we drop in the embed code, the job list will use MezTal's typography and color scheme, maintaining consistency. We will verify this and then apply additional CSS tweaks if needed. Workable provides special CSS hook classes (e.g. `.whr-title`, `.whr-location`) that we can target ⁸ ⁹. For example, if our site uses a dark background in the footer where jobs are embedded, we'll ensure the text from Workable appears in a contrasting color (we might need to override the default if inheritance isn't perfect). We will hide any fields we don't need (like job codes or post dates) via CSS as demonstrated in their docs ⁹ to keep the design clean. In essence, we treat the Workable widget as an extension of our site: same spacing, responsive behavior, etc. The widget is mobile-responsive out of the box, but we'll test it. If job descriptions are embedded, we might limit their height or use an accordion style so the page isn't overwhelmingly long – Workable allows showing full descriptions or just titles. We'll likely show a short list of open roles (title, location) with a "View more" link to Workable's page if needed, to conserve space and style.
- **LinkedIn Content (Posts or Feeds):** To showcase LinkedIn content (like our latest posts or testimonials), we have a couple of options. LinkedIn offers an official embed for individual posts. For example, we might embed a specific post (like Gemini's thought leadership post) using LinkedIn's provided iframe script. This embed will carry LinkedIn's branding (white background, LinkedIn logo, etc.) which might clash with our design slightly. Our plan is to **contain LinkedIn embeds within a styled container** that matches our site – e.g., place it on a white section of our page if the embed is white, or use CSS to give the iframe a drop shadow or border that aligns with our style guide. We cannot fully restyle LinkedIn's internal widget (since it's an iframe), but we can control its placement and surrounding elements. Another approach is using a third-party LinkedIn feed widget (like Elfsight or SociableKIT) which allows more customization of the feed's appearance. If we need to show a feed of posts, we might use such a tool to pull the posts and match them to our style (fonts, colors). The key is to "**match our branding**" as much as possible ¹⁰ – for instance, choosing a post layout that uses neutral colors that complement our palette, and ensuring any clickable elements use our accent color if possible.
- **Consistent Framing:** We will introduce embedded content with proper headings or frames so it looks intentional. For example, above the LinkedIn feed embed, we might have a section header "Latest on LinkedIn" styled like our other section headers. This prepares the user for the content and also frames the widget aesthetically. Similarly, for Workable, if we embed jobs on a page that's not solely the careers page (say the homepage has a short "We're hiring" widget), we'll put it in a card or box styled with our brand colors.
- **Responsive Behavior:** We will test the embedded elements on various screen sizes. The Workable widget generates a list that should stack on mobile. If any part overflows, we'll add custom CSS

(within `<style>` tags as Workable suggests) to adjust font sizes or spacing for small screens. LinkedIn post iframes have fixed widths; we'll ensure to wrap them in a container with `max-width: 100%` so they shrink on mobile. If using a third-party LinkedIn widget, we'll configure it to display a single-column feed on mobile.

- **Loading and Performance:** Embeds can sometimes slow down a page. We'll mitigate this by **lazy-loading** these sections if possible. For example, we could load the LinkedIn script only when that section scrolls into view (using Wix's Velo code or a third-party plugin). For Workable, we'll integrate it on the Careers page primarily, so users going there expect to load job info. We'll keep an eye on not embedding too many heavy external scripts on critical pages.

In summary, **Workable's jobs widget will seamlessly adopt our site's style** (verified via their docs: it inherits CSS ⁷, and additional tweaks are possible), and **LinkedIn content will be embedded in a way that complements our design**, either via careful use of official embeds or using a customizable feed tool. Our design team will ensure fonts, colors, and spacing of these embeds feel native to the page. For instance, any headings within the Workable widget can be adjusted to our heading style (via CSS overrides), and the LinkedIn feed's post style can be chosen to match our branding ¹⁰. By doing this, even though the content comes from external sources, the user experience remains cohesive and professional, reinforcing the MezTal brand rather than looking like pasted-in widgets.

(Sources: Workable's documentation confirms embedded widgets inherit site CSS, making styling easier ⁷. Also, best practices for embedded social feeds suggest matching the feed style to branding ¹⁰.)

6. Schema Markup Decision for Dynamic Role Pages

(`/hire/[slug]`)

To prevent Google from misclassifying our role pages as job postings, we will **avoid using Job Posting schema** on these pages and instead use an appropriate schema that reflects their true purpose. The dynamic "Hire [Role]" pages are essentially **service/solution pages describing our ability to staff that role**, not actual job listings for applicants. Therefore, applying `JobPosting` structured data would be incorrect and could mislead search engines (and might trigger inclusion in Google Jobs, which we do not want).

Chosen Schema: We will mark up these pages as **Organization > Service** pages (Schema.org *Service* type), possibly with an additional type of *WebPage* (or *Article*) for completeness. Using *Service* schema is recommended for pages that describe offerings of a business ¹¹. In our case, each role page is describing a staffing service for that role (for example, "Provision of Senior Living Billing Analysts" is a service we offer). This tells Google the page is about a professional service, not a job vacancy. We'll include properties like service type (the role), area served (Latin America/U.S.), provider (MezTal as Organization), etc. This approach aligns with general SEO guidelines for agencies – service pages should use Service schema ¹¹. We will also include the standard *Organization* schema (site-wide in the footer or JSON-LD) to identify MezTal as the provider, and link the Service to the Organization via the "provider" or "serviceOperator" field.

Additionally, each role page often contains rich content (FAQs, potentially a blog-like narrative). We might also apply `FAQPage` schema for the FAQ section if present, and/or mark the main text as an `Article` or `WebPage` (with `WebPage@pageType: "WebPage"` and a suitable name). However, we must **not** include

`JobPosting` schema. By clearly avoiding any JobPosting markup, we signal to Google that these pages are not part of a jobs listing board.

We'll double-check that nothing on the page triggers Google's automated Job Posting detection. For instance, sometimes Google might algorithmically think a page with "Salary: \$X" and "Requirements:" looks like a job ad. To counter that: - We will contextually phrase content to emphasize it's about **hiring service**. E.g., instead of a plain list of "Job Requirements" we might say "**Common Skills** our *[Role]* candidates have" – framing it as informational, not an actual job requirement list. - We might include a blurb like "*MezTal provides experienced [Role] professionals***" etc., so it's clear it's an offer to hire out talent, not hire in-house.

Using the Service schema will further solidify this. We'll include `"name": "Hire [Role] Service"` in the markup, and possibly `"category": "Staff Augmentation"` or similar. This approach is recommended in SEO circles for agencies to differentiate service pages from other content ¹¹. In fact, recruitment agencies are often advised to use Service or Organization schema on their role pages and only use JobPosting on actual job ads meant for candidates.

Outcome: With this schema strategy, Google's crawlers will correctly interpret our `/hire/[slug]` pages as business service pages. This should keep them out of Google Jobs results (which require JobPosting markup or clear signals of a job listing). Instead, they can appear as normal search results, potentially even earning rich snippets for FAQs if we include that markup. We'll also leverage the schema to enhance E-E-A-T: including the Organization schema ties the content back to MezTal (establishing who we are), and perhaps Author schema if we have someone (like Gemini) "authoring" these pages for credibility.

To implement, we will add JSON-LD in the page template. Wix allows page-level JSON-LD injection via the SEO settings for dynamic pages. We will programmatically include each role's data (role name, etc.) in the JSON. For example:

```
{
  "@context": "https://schema.org",
  "@type": "Service",
  "name": "Hire {{RoleName}}",
  "description": "{{someBriefDescription}}",
  "provider": {
    "@type": "Organization",
    "name": "MezTal",
    "url": "https://www.meztal.com"
  },
  "areaServed": "Americas",
  "serviceType": "Nearshore Staffing"
}
```

(the above is illustrative – we'll refine properties based on Schema.org specs). By doing this, we not only avoid the JobPosting trap, we also make the page eligible for any rich results applicable to Service (and at minimum, we're doing proper structured data for SEO hygiene).

In summary, we will not use `JobPosting` markup on role pages (since no specific job is being offered to applicants), we'll use `Service` schema (and possibly WebPage/FAQ as secondary types) to clearly label the content's nature ¹¹. This ensures Google treats the pages as service descriptions in the staffing domain, not as job ads – preventing any confusion or misplacement in search features.

(Reference: Schema guidance for agencies suggests using Product/Service schema for service pages instead of job posting ¹¹, which aligns with our approach.)

7. Lead Form Attribution & Routing for CMS-Based Dynamic Pages

Every dynamic page (whether it's a role page, location page, or other CMS-driven page) will include a **lead capture form**, and we need to attribute form submissions to the page of origin so that leads are correctly routed internally. Our plan involves two key steps: **capturing the page context** with each form submission, and **using that data to route the lead in our CRM (HubSpot)**.

Technical Implementation (Attribution): We will add a hidden field to the form, e.g. “**Source Page**” or “Page URI”. In Wix (or whichever form solution we use, possibly HubSpot embedded form), we can inject the current page’s identifier into this field. For instance, on a dynamic role page, we can use the page’s *role name* as the hidden value. Wix Code can set the field value on form load: `$w("#sourceField").value = dataset.getCurrentItem().title` (if using a dataset) or simply grab `window.location.href` if needed. If using a HubSpot form embed, HubSpot can automatically capture the page URL and page title as metadata; we will double-check this and map those to properties. Regardless, we will end up with a submission property like “Source Page: Hire-Python-Developer” or a URL containing “/hire/python-developer”.

CRM Routing Logic: In HubSpot (or our marketing automation tool), we’ll set up workflows that look at this “Source Page” field (or the page URL). The logic will route or tag leads based on the content. For example: - If Source Page contains “senior-living” or the role is a healthcare one, we might assign the lead to our sales rep specialized in Healthcare accounts. - If the form came from a Location page (say “Guadalajara”), perhaps it routes to the team handling nearshore center inquiries. - If it’s a generic blog or resources page, it might go to a general queue.

At minimum, we will attach the page info to the contact in HubSpot so that our sales team can see context: “*Lead downloaded form from Hire Python Developer page.*” This is invaluable for lead nurturing – the rep immediately knows the lead’s interest area (in this case, Python developers) and can tailor the follow-up. HubSpot allows dynamic workflows, so we’ll configure something like: - **Trigger:** Form submitted on *Hire Role* pages. - **Action:** Assign HubSpot lifecycle stage “Marketing Qualified Lead – Role: [X]” and create a task for sales owner Y (if we have specialized owners).

Our HubSpot closed-won data indicates that leads often come in with specific needs (e.g., one client’s first inquiry was specifically for a “Billing Analyst” role, which turned into a deal). Capturing that exact need via the form’s page context is crucial. We saw in the data that Trustwell Living came through seeking a Billing Analyst and later closed – if we hadn’t known they wanted a Billing Analyst, we might not have paired them with the right specialist. With our attribution setup, **every form submission brings its context along**.

On the Wix side, implementing the hidden field is straightforward. We'll test a few submissions to ensure the data flows into HubSpot. We will also incorporate UTM parameters if campaigns are driving traffic, but that's separate from page context (those handle external source, while our hidden field handles on-page context).

Form Variations: We may have slightly different forms on different page types (for example, a role page might have a form titled "Hire this talent" vs. a location page form "Find talent in [Location]"). In the CMS, we can even store a form ID or short code if needed and dynamically load it. But likely one general form with dynamic field values will do. The key is that the user doesn't have to manually specify what they're interested in – the system knows based on where they submitted.

Routing Examples: - A lead from a "*Hire React Developer*" page might be tagged in HubSpot with **Interest: React Developer**, and our SDR team's workflow could automatically send them our React developer hiring guide email, and assign to the tech industry SDR. - A lead from the *Guadalajara* page could be tagged **Interest: Guadalajara hiring**. Internally, we might route that to a representative of our Mexico team. Also, in HubSpot, we could increment a counter for interest in Guadalajara to inform our marketing (e.g., "X leads in last quarter showed interest in GDL").

To ensure no lead falls through cracks, we'll set a default: if no specific routing rule matches, leads go to a general pool, but with the page info in the notes. In HubSpot, this might simply create a ticket that includes the page URL and we triage manually.

Data Management: We'll keep an eye on that "Source Page" taxonomy. For consistency, we might use the page's slug or a friendly name. Since these are dynamic pages, the slug is unique and likely self-explanatory (e.g., "hire-python-developer"). We'll likely use the page title (which might be "Hire Python Developer – Nearshore IT Staffing") to populate the field. That clarity will carry into HubSpot. (We'll confirm that long titles don't break anything – HubSpot can handle a decent length text field.)

In conclusion, **each dynamic page's form will silently capture the page identity and pass it through**, enabling us to auto-segment and route leads. This ensures, for instance, that a niche role inquiry (like *Clinical Recruiter*) is recognized and perhaps flagged for a salesperson who specializes in clinical staffing. It validates our new site's logic by aligning our sales process with the CMS structure: whatever content attracted the lead, our follow-up will continue the conversation on that topic. This tight coupling of content and CRM is a big improvement — no more guessing what a contact is interested in; we'll know from the moment they convert.

(We did not find obstacles in current data for this; our HubSpot integration supports capturing page URL, and our closed-won logs underscore the importance of knowing the role of interest when a lead comes in.)

8. Special Content to Elevate Guadalajara as the Canonical Geo Hub

Among our *Locations* CMS entries, **Guadalajara** will be treated as the flagship page – effectively the pillar of our geographic hub content. To elevate it as the “canonical” geo hub, we’ll take several steps in content and linking strategy:

- **Comprehensive Content:** The Guadalajara page will be more extensive than other city pages. While other city pages (e.g., Puebla, Mérida) might have ~500 words of content, Guadalajara’s page will be a robust guide (potentially 1500+ words) about why Guadalajara is the prime location for tech talent in Mexico. It will include sections on the local tech ecosystem, universities, major companies present, talent pool statistics, and maybe cost-of-living advantages. By packing it with valuable info (and perhaps a custom infographic or two), it becomes the **authoritative resource** on nearshore staffing in Guadalajara. This not only serves SEO (targeting keywords like “Guadalajara tech talent” etc.) but also positions it as the reference page for anything related to LatAm hubs.
- **Internal Linking Favoritism:** We will link to the Guadalajara page from multiple places: the main navigation (e.g., in a “Locations” dropdown, Guadalajara might be listed first or even separately highlighted), from the homepage (maybe a callout like “Mexico’s Silicon Valley: Why Guadalajara Leads in Tech Talent”), and from other location pages. On every other city’s page, we can include a subtle prompt like: *“Looking for the largest tech hub? Check out Guadalajara – our central hub for IT talent.”* This cross-link signals that Guadalajara is the central node. Essentially, we create a hub-and-spoke where **Guadalajara is the hub** and other cities are spokes linking back to it. This also helps with SEO by funneling authority to Guadalajara’s page. Intugo’s strategy illustrates this concept: they physically highlight presence in major hubs like Guadalajara to build credibility and rank for city keywords. We’ll mirror that by ensuring Guadalajara is prominently cited and linked across the site.
- **Canonical URL / SEO Canonical Tag:** If there are overlapping topics (for example, maybe we have a general “Mexico IT Talent” page or blog posts discussing Mexican cities), we might set the canonical URL for broad Mexico-related content to point to the Guadalajara page. This is more of a technical SEO tweak: for instance, if an article “Top 5 Cities in Mexico for Developers” covers similar ground, we ensure proper canonical tags or at least link that article strongly to Guadalajara’s page. However, since Guadalajara’s page is unique, the main “canonical” notion here is metaphorical (i.e., the primary content source on geo topics).
- **Rich Media and Showcasing:** We can embed a video or slideshow on the Guadalajara page – perhaps a short video of our Guadalajara office or tech community events. Also, success stories specifically from Guadalajara (client case studies where the team was centered there) could be featured. This gives the page human interest and weight that others might not have. It subtly tells Google and users “this is the main hub page”.
- **Schema Markup:** We might use `LocalBusiness` or `Place` schema for Guadalajara since it’s a physical hub. For example, we could mark it up with `"address": Guadalajara, JAL, MX` and mention it as a headquarters. However, since MezTal’s HQ is officially in the US, we must be careful not to confuse. Instead, we might use `City` or just ensure Organization schema mentions a Guadalajara location. The point is to reinforce through structured data that Guadalajara is a key location for us. This, combined with content, helps search engines associate MezTal strongly with Guadalajara.

- **Authority signals:** We'll include external validation where possible – e.g., mention partnerships with Guadalajara tech parks, or quote a stat like "Guadalajara produces 8,000 engineering grads per year" with a citation. If we have a quote from Gemini or our CEO about Guadalajara's importance ("We chose Guadalajara for our Delivery Center because..."), we'll put that on the page. These elements not only add unique value (making the page link-worthy for others) but also solidify its role as the go-to page for all things nearshore-Mexico.
- **Navigation structure:** In the sitemap, Guadalajara will likely live under "/locations/guadalajara" as others do. But we might also have a top-level page about "Nearshore IT Staffing" under which Guadalajara is a prominent link (in fact, our SEO plan shows *Nearshore IT Staffing* as a pillar and *Guadalajara* as a key page under geographic hubs). We will ensure no two pages compete for the same keyword. For example, if we have a generic "Mexico IT Outsourcing" page, we'll either combine it with Guadalajara or have it heavily reference Guadalajara as the main content.

In effect, we treat the Guadalajara page as the **pillar page for our "Geographic Hubs" content cluster**. All other city pages (Monterrey, etc.) will link up to it (through nav or in content), and it will link out to them in a controlled way (maybe a section "Other Locations in Mexico" listing the others). But Guadalajara will always be portrayed as the flagship: more content, more links, more prominence.

This approach is validated by competitor moves: Intugo emphasizes its physical presence in GDL and other major hubs as a trust signal. Similarly, we call Guadalajara our "center of excellence" or similar phrasing. We might even include an element on the homepage that says "**Nearshore Center: Guadalajara, Mexico**" with a snippet of that content, clicking through to the full page.

Why Guadalajara? It's worth noting in content that Guadalajara is often called the "Silicon Valley of Mexico." We will use that narrative. Possibly, we'll maintain a blog or updates on that page (like a section listing upcoming tech events in Guadalajara) to keep it fresh.

By doing all the above, when Google looks at our "locations" section, it should easily identify Guadalajara's page as the most important (lots of internal links pointing to it, richest content). Users too will notice – if they explore locations, they'll see one page that clearly stands out, funnelling them to it. From there, they can branch out to others if needed, but we've successfully communicated that **Guadalajara is our key hub** (which aligns with our operations and strategy).

Finally, as a canonical hub, the Guadalajara page can serve as a **landing page for campaigns**. Any general "Hire in Mexico" AdWords or LinkedIn ad we run might land on that page. We'll ensure the content is conversion-optimized (strong CTA for getting talent in Guadalajara). In contrast, smaller location pages might not be used for ads but purely organic support.

In summary, through **content depth, strategic linking, and emphasis in site architecture**, we will elevate Guadalajara's CMS record to be the undisputed central geo page. It will be the page that both users and search engines associate most strongly with "nearshore IT talent in Mexico," and other location pages will play supporting roles to it.

(Competitor precedent: Intugo reinforces credibility by highlighting presence in Guadalajara, among other hubs, which supports our strategy to make GDL highly authoritative.)

9. Redirect Logic from Legacy Sitemap to New URLs (e.g., /hire/python → /nearshore-it-staffing/hire-python)

When we launch the new site with its Hub-and-Spoke URL structure, it's critical to put 301 redirects in place mapping all old URLs to their new counterparts. Our legacy site had a "flattened" URL scheme (for example, role pages directly under /hire/ at root, and perhaps other pages like /services or blogs without hierarchical folders). The new site organizes content into folders like /nearshore-it-staffing/hire-python, which is more SEO-friendly but different. We will implement a **comprehensive redirect plan** to ensure any traffic or SEO equity from the old links flows to the correct new pages.

Identification of Legacy URLs: First, we've compiled the old sitemap (from meztal_sitemap.xml and Google Analytics data). Notably: - Old role pages were at /hire/[role] (e.g., /hire/python, /hire/react, perhaps /hire-senior-living-billing-analyst if it existed or simply not at all if we didn't have it). - Some old pages might have been at the root or other structures (for example, comparisons could have been like /staff-aug-vs-outsourcing etc., and now they live under /comparison/...). - Blog posts likely were under /blog/[slug] and remain similar, but we should confirm if the new blog uses a different slug pattern (maybe /resources/blog/[slug] now). - Location pages might be entirely new (the old site may not have had dedicated city pages, so that's new content – but if any old pages discussed locations, we redirect those to the new location pages or the Guadalajara page as appropriate).

Redirect Mapping: We will create a mapping table, for example:

/hire/python	->	/nearshore-it-staffing/hire-python
/hire/react	->	/nearshore-it-staffing/hire-react-developer (assumed new slug includes "-developer" perhaps)
/hire/java	->	/nearshore-it-staffing/hire-java-developers

Basically, any /hire/[tech] goes to /nearshore-it-staffing/hire-[tech] (with slight slug adjustments if we changed pluralization or wording). This can be done with a pattern redirect: /hire/:slug → /nearshore-it-staffing/hire-:slug (assuming slug formats match). In Wix, we might have to enter individually if pattern isn't supported, but we can certainly bulk upload these via the SEO tools.

For other sections: - If the old site had URLs like /nearshore-it-staffing (maybe it didn't, because that's new), we ensure none conflict. - Old comparison pages (if any, e.g., maybe /toptal-vs-meztal previously) now may live at /comparison/meztal-vs-toptal (just hypothetical). Those need redirects too. - We must also handle **case insensitivity** and trailing slashes if any (usually not an issue if consistent).

Importantly, some old URLs might not have direct one-to-one new pages. For example, if /hire/python existed but now we have it under a new parent, that's straightforward. But what about an old page like /hire-python (no subfolder)? In the prompt, they specifically mention "flattened URLs like /hire/python instead of /nearshore-it-staffing/hire-python." It implies some old URLs might have actually been in a combined form like /hire-python (with a hyphen). We need to clarify: - Possibly the old site had /hire-python (with no directory, just as a page under root) because "Hire Python" might have been seen as one

concept. And now we have `/nearshore-it-staffing/hire-python`. - If so, we will redirect `/hire-python` → `/nearshore-it-staffing/hire-python`. Similarly `/hire-react-developers` → `/nearshore-it-staffing/hire-react-developers`.

We will use a mixture of **exact redirects and pattern-based** where safe. For instance, anything starting with `/hire-` at root can redirect to `/nearshore-it-staffing/hire-...` (same slug after hire-). We just must be cautious that we don't inadvertently redirect something unintended.

Testing & SEO Considerations: All redirects will be **301 (permanent)** so that search engines pass the link juice. We'll test key ones manually (type old URL, see it land on the new). Wix allows managing 301s via its dashboard; we will input them there. If any old pages are irrelevant or we chose not to carry them over, we still redirect them to a close relevant page (never let them just 404). For example, if an old page was a very thin page about "hire software developers" and we now have a richer "hire a developer" page, we map old to new.

We'll pay attention to query parameters too – if some links had `?source=linkedin` etc., Wix will by default carry them over on redirect. That's fine.

Legacy Sitemap Handling: Once redirects are in place, we'll update the robots/sitemap such that only new URLs are listed. The legacy sitemap's URLs will either be gone or pointing to new pages. Anyone hitting an old URL from search results will be seamlessly taken to the new page that most closely matches their intent.

Example given (/hire/python): This will redirect to `/nearshore-it-staffing/hire-python` [69t] (the new intended location for Python developer hiring content). Another example: an old blog like `/blog/remote-vs-onsite` might now live under `/resources/blog/remote-vs-onsite` – we would redirect accordingly.

Where we changed the slug naming, we'll account for that. E.g., if "Staff Augmentation" page was `/staff-augmentation` and now it's `/nearshore-it-staffing/it-team-augmentation`, we map the old to the new even though the slug changed more significantly. Each such change has been catalogued during our content migration planning.

Monitoring: Post-launch, we'll monitor Google Search Console for 404s or redirect chains. If we see any missed ones (perhaps an obscure old URL), we'll add a redirect for those too. The aim is 100% continuity.

In sum, **every URL from the legacy site will either have an exact new location or be pointed to the closest relevant new page**. Our redirect logic ensures no SEO value is lost and users using old bookmarks or search results are taken to the right content. The specific mention of flattened `/hire/python` to new nested URL is handled via a rule covering all `/hire/...` pages. We'll similarly handle any other "flattened" patterns. This is a one-time setup but extremely crucial for SEO preservation during the site rebuild.

(We did not find a direct snippet in provided data about redirects, as it's a straightforward SEO task. However, this plan is derived from standard best practices for site URL changes.)

10. Distinction Between Blog and Resources Collections (Templates & Usage)

We have two separate CMS collections for content marketing: **Blog** and **Resources**. It's important to clarify their differences in purpose, content format, and how their page templates will function, since this impacts how we create and present content in each.

Purpose & Content Type: - **Blog:** The Blog collection will contain regular blog posts – typically timely or thought-leadership content, news updates, opinion pieces, and SEO-driven articles that are more *editorial*. These are usually informal or narrative in style, often targeting top-of-funnel queries. For example, “5 Tips for Managing Remote Teams” or announcements like “MezTal Opens New Office in Guadalajara” would be blog posts. Frequency is higher, and each post has an author and date. Blogs are meant to engage readers with insights and encourage sharing/comments. - **Resources:** The Resources collection is for more **evergreen, long-form content** and assets that prospects might consult as references. This includes things like in-depth guides, whitepapers, case studies, e-books, possibly webinars or downloadable content. These pieces are often mid-to-bottom funnel. For example, “Nearshore Staffing Ultimate Guide 2024” or “Case Study: How X Company Scaled with MezTal” would live in Resources. Resources content tends to be more polished and evergreen (updated occasionally but not tied to a specific publish date for relevance).

In our site architecture, we envisioned “Resources” as a broad pillar that might encompass various content types ¹². Practically, we decided to implement it as a separate collection from the Blog because: - We likely want a **different template layout** for resource pages (more like a library entry, possibly with a sidebar for download or a summary box). - Resources might have additional fields that blogs don't: e.g., a PDF download link, an “Excerpt” or “Executive Summary”, or fields like Industry (for case studies) or Type (Guide, Case Study, Checklist). - Conversely, Blogs have fields like Author, Featured Image, Published Date, which Resources might not display prominently (a case study might not list an author or exact date on the page, or might list a date in a subtle way).

Template Requirements: - **Blog Template:** The blog post page will show the title, author name and bio, date, and content (rich text). It might have a comments section or social share buttons. It will also likely show categories or tags for navigation to related posts. The design will be optimized for reading (likely a single column of text). Because blogs are chronological, we'll have a blog listing page that paginates by date (newest first). That listing will pull from the Blog collection only. Also, the blog template might include a sidebar with perhaps a search or recent posts list. - **Resources Template:** A resource page might be more structured. For example, a case study page might have a **hero section with a summary** (client name, results achieved, etc.), then a breakdown of challenges/solutions. A downloadable guide page might have a **prominent CTA to download the PDF** (“Download PDF” button) possibly gated by a form. The template will accommodate these. We may even have sub-templates within Resources depending on type (if case studies look very different from guides, we could either differentiate by a field layout or use two collections, but currently we plan one Resources collection with a field to distinguish type). The Resources template likely won't show author or casual blog meta info. Instead, it might show something like “Last Updated on [Date]” if relevant (for guides we update). It might also cross-link to related resources by category.

Think of it this way: **Blog is a subset of Resources in a broad sense (both are content)**, but we treat Resources as a *premium content library*. That's why on the site we might label navigation as “Resources” which includes blog posts and other content ¹². But under the hood, we separate the collections to

manage them better. The blog posts flow into the regular cadence and RSS feed, whereas resources might be fewer and more static.

User Experience Difference: If a user clicks “Blog”, they get an index of articles (with dates). If they click “Resources”, we might present a different index – possibly a page that filters by resource type: e.g., tabs for Guides, Case Studies, etc. We could actually build the Resources landing by pulling items from the Resources CMS. Perhaps we’ll have tags or a category field for type of resource, and allow filtering. This is easier to manage with a separate collection. Meanwhile, the Blog landing is just the blog roll.

Also, resources might integrate with marketing workflows (like require form fill to access certain content). For example, a high-value whitepaper could be a Resource entry where the “Download” button triggers a HubSpot form popup. The blog posts, on the other hand, are freely accessible and ungated always.

SEO and Metadata: Blog posts will use `BlogPosting` schema (a subtype of Article) which fits their nature (and can get date-rich snippets, etc.). Resource pages might use `Article` or some specific schema (CaseStudy schema exists, for instance, for case studies – we might use that for those entries). We’ll tailor metadata accordingly. For instance, blog posts will have meta descriptions that might include the date or be more catchy, while resources might have more formal meta descriptions focusing on the value proposition.

Why Not One Collection? The decision to have two is to avoid clutter and complexity. We expect a larger volume of blog posts vs. fewer resource pieces. If in one collection, mixing them could complicate querying (e.g., showing only “guides” vs “posts”). Also, content managers in our team might be different – a content marketer might handle blogs, whereas a strategist handles case studies. Separating collections gives us workflow control (different permissions or just easier mental separation).

Template Example Differences: - A **Blog post page** might show: “**How to Improve Agile Communication**” by Jane Doe, October 10, 2025. Then the content, then maybe “About the author” box, and “Related Posts”. - A **Resource page** for a case study might show: “**Case Study: Scaling FinTech Team in Guadalajara**” (no author or date at top, perhaps a subtitle “Client: FinCorp – 12 Engineers hired in 6 months”). It might have a sidebar with key metrics (e.g., timeline, team size, results) and a PDF download. At the end, maybe a CTA “Contact us to learn more”. Related resources might be other case studies or guides, not blog posts.

From a maintenance perspective, we’ll ensure the CMS fields reflect these needs. The Blog collection likely has fields: Title, Slug, Author (ref to Authors), Date, Featured Image, Content, Category/Tags. The Resources collection may have: Title, Slug, Type (dropdown: Guide/Case Study/etc.), Thumbnail, Short Description, Content (rich text or multiple fields sections), perhaps File (for downloadable) or Video link, etc.

Finally, on the front-end, we’ll distinguish them in navigation labels to avoid confusion. Possibly the main menu has “Resources” and under it “Blog” separately, or vice versa. We might even rename one in the nav to be crystal clear (some sites use “Insights” for blog and “Resources” for downloads). But given the prompt, likely “Blog” and “Resources” are distinct sections.

In conclusion, **Blog and Resources are distinct both in content strategy and template design:** - *Blog = frequent, topical articles managed in Blog CMS, standard blog layout.* - *Resources = high-value content pieces in Resource CMS, more structured page layout per piece.*

This separation ensures each content type can shine with the appropriate format. It also helps SEO: blog posts feed the ongoing content machine (good for regular indexing and broad topics), while resources act as evergreen cornerstones that our sales team and inbound marketing can leverage (often aligning with bottom-funnel keywords or offering in-depth knowledge). Our site architecture, as noted in previous planning, treats “Resources” as a broad category encompassing multiple content kinds ¹² – by engineering it with two collections, we achieve that breadth without one template trying to do it all.

(Context: In the site outline we prepared, “Resources” was a pillar including blogs and guides ¹². The decision to use two CMS collections is a practical implementation choice to meet different template requirements.)

Next Steps Checklist

- **Finalize CMS Data & Taxonomy:** Complete the mapping of all roles to locations (reference fields) and verify Salary_Range entries for each role (identify data sources for any **DATA GAP** on salaries and populate accordingly).
- **Implement Schema Markup:** Add JSON-LD templates for Service schema on role pages, Organization info site-wide, and appropriate Article/BlogPosting schema on Blog posts and Resource pages (use CaseStudy schema for case studies, etc.). Test with Google’s Rich Results tool.
- **Configure Redirects:** Prepare and upload the 301 redirect list from all legacy URLs to new URLs. Double-check critical ones like `/hire/[tech]` → `/nearshore-it-staffing/hire-[tech]`. After launch, monitor for any 404s and add redirects as needed.
- **Design QA for Embedded Content:** Embed the Workable widget and a sample LinkedIn post on a staging page to ensure they inherit styling. Apply custom CSS per Workable’s docs (e.g., hide unwanted fields, adjust fonts) ⁷. Ensure LinkedIn embeds are responsive and match branding (consider using a third-party feed if the official embed looks off-brand) ¹⁰.
- **Content Enrichment for Guadalajara:** Develop the extended content for the Guadalajara page (gather stats, write copy, design infographic). Likewise, draft slightly lighter content for other locations. Cross-link all location pages appropriately, with Guadalajara as the central hub. Internally link from various pages (homepage, etc.) to Guadalajara’s page to signal its importance.
- **Form Testing & CRM Automation:** Set up the hidden “Source Page” field on forms and do test submissions from a role page and a location page. Verify in HubSpot that the data comes through (page title or slug captured). Then create HubSpot workflows to assign leads based on that field (e.g., if contains “Guadalajara” or a specific role). Ingrain a process for sales to receive notifications including that context.
- **Template Differentiation Review:** Review the Blog vs Resource templates side by side. Ensure Blog posts show author info and dates clearly, and Resources pages have the necessary download or structured sections. Check that the Resources listing page cleanly segments content (e.g., filter by type or at least visually distinguish case studies vs guides via labels).
- **Quality Assurance & Launch Prep:** Before launch, crawl the new site to ensure all internal links point to new URLs (no self-inflicted 404s). Check that schema markup is not flagging errors (especially make sure no JobPosting schema slipped in). Also, verify that key pages (role pages, Guadalajara page) are included in the XML sitemap with correct updated URLs.
- **Post-Launch Monitoring:** After going live, monitor Google Search Console for any coverage issues or misinterpretation (for instance, check that role pages are not listed under Google Jobs – they should appear as normal web results). Monitor analytics to see if leads are indeed filling forms on the new pages and that our attribution is capturing as expected.

By following this checklist, we'll validate that the strategic execution (content, technical SEO, and CRM integration) is working as planned, and we'll be able to catch any issues early and iterate. Our new Hub & Spoke model in Wix will then be fully operational, with each piece (roles, locations, comparisons, blog, resources) supporting the others in a cohesive ecosystem.

1 How AI Is Redefining How Recruitment Agencies Appear in Search - Kaizen SEO

<https://www.kaizen-seo.com/blog/how-ai-is-redefining-how-recruitment-agencies-appear-in-search-169>

2 3 4 5 6 Near vs. Toptal: Which Is Better for Hiring Remote Talent?

<https://www.hirewithnear.com/compare/near-vs-toptal>

7 8 9 Can I change the styling options for the embedded job widget? – Workable Help

<https://help.workable.com/hc/en-us/articles/115012807127-Can-I-change-the-styling-options-for-the-embedded-job-widget>

10 LinkedIn Widget — Embed LinkedIn Feed on your website

<https://elfsight.com/linkedin-feed-widget/>

11 Schema Markup – Is It Really Optional for Agencies? – The Admin Bar

<https://theadminbar.com/seo-weekly/schema-markup-is-it-really-optional-for-agencies/>

12 meztal_chat_summary.txt

<file:///file-HPkVqDhmpD6HWcbDRxazqV>